

1) High-level design and decisions (what to choose and why)

1. Purpose & scope — decide what you want to learn:
 - Low-interaction (e.g., port responders) vs high-interaction (full SSH shell) — low interaction is safer and cheaper; high interaction gives richer data (real attacker activity and payloads).
 - Single honeypot vs honeynet (many sensors over regions) — start single, then scale.
 2. AWS building blocks you'll use:
 - EC2 (honeypot instances), Security Groups, VPCs/subnets (isolation), IAM role for logging, S3 for storing logs, VPC Flow Logs, CloudWatch, and optionally GuardDuty for cross-analysis/alerting. AWS guidance recommends VPC security best practices and using GuardDuty and flow logs as part of monitoring. [AWS Documentation+1](#)
 3. Choice of honeypot platform:
 - Cowrie — SSH/Telnet interactive honeypot, well-documented and lightweight (good single EC2 sensor). [codygula.com](#)
 - T-Pot — all-in-one platform (many honeypots + ELK) useful for single VM full stack. [InfoSec Write-ups](#)
 - MHN (Modern Honey Network) — server + distributed sensors to centrally manage many honeypots; good when scaling. [GitHub](#)
-

2) Architecture (recommended for production-ish testing)

- Create a dedicated AWS account or separate VPC (do NOT put honeypots in production VPC).
- VPC with public subnet for honeypot instances (internet reachable) and a separate management subnet (private) for logging/management servers (ELK / MHN server).
- Security Group for honeypot:
 - Inbound: open specific service ports you want to attract (e.g., TCP 22 and TCP 80) to 0.0.0.0/0 **only** for honeypot ports.
 - Inbound: admin SSH port (e.g., 2222) only from your IP.
 - Outbound: restrict where possible — but note some honeypots require outbound access to fetch updates; if you want to prevent pivoting, block outbound except to necessary services.

- VPC Flow Logs → S3 (and optionally Athena) for network forensic queries.
 - Central logging: send honeypot logs to S3 and also forward to ELK/OpenSearch or to MHN/T-Pot UI.
 - Enable GuardDuty and CloudTrail for cross-correlation/alerts. [aws.github.io+1](https://aws.github.io/t-pot/)
-

3) Concrete step-by-step: deploy Cowrie on AWS EC2 (single honeypot)

This is a compact, tested pattern — Cowrie on Ubuntu EC2, logs to S3 + optional ELK.

Preliminaries (assumes AWS console + CLI available)

- Create/identify an AWS account and region.
- Create an SSH key pair in the region (or import one).
- Note your admin public IP (for locking admin access).

Step A — Create VPC & subnets (isolation)

1. In AWS Console → VPC → Create VPC (CIDR: 10.10.0.0/16).
2. Create a public subnet (10.10.1.0/24) for honeypot.
3. Create an internet gateway and attach it.
4. Create a route table for the public subnet with 0.0.0.0/0 → IGW.

(You can also use the default VPC but dedicated VPC is safer.)

Step B — Create Security Group

- Name: honeypot-sg
- Inbound rules:
 - SSH for admin: TCP 2222 — Source: your-ip/32
 - Cowrie (fake SSH): TCP 22 — Source: 0.0.0.0/0 (if you want to attract internet scanners)
 - HTTP (if running web honeypot): TCP 80 — 0.0.0.0/0
- Outbound rules:
 - Deny all by default, then allow minimal egress (e.g., DNS/HTTP) if you need updates. Consider using a NAT for controlled egress to a proxy.

Important: Ensure your real admin SSH listens on a non-default port (2222) and is restricted to your IP. Cowrie will bind to port 22 for the fake shell. Several tutorials recommend this pattern. codygula.com

Step C — Launch EC2 (Ubuntu 22.04 recommended)

- Instance type: t3.micro or t3.small (for low traffic). Choose larger if expecting heavy activity.
- AMI: Ubuntu Server 22.04 LTS.
- Subnet: public subnet created earlier.
- Assign Public IP: yes.
- Security group: honeypot-sg.
- Key pair: your key.

Step D — Basic hardening for management

SSH to instance using admin port (example if you open port 2222 for admin):

```
# from your workstation (adjust username/ec2-user accordingly)  
ssh -i mykey.pem -p 2222 ubuntu@<EC2_PUBLIC_IP>
```

On the instance:

```
# update & install essentials  
  
sudo apt update && sudo apt upgrade -y  
  
sudo apt install -y git python3-venv python3-pip build-essential libssl-dev libffi-dev  
  
# create a non-root user for cowrie if needed
```

Step E — Install Cowrie (quick install)

This example follows the common approach used in several guides — configure a virtualenv and run Cowrie as a service.

```
# create directory  
  
sudo adduser --disabled-password --gecos "" cowrie  
  
sudo su - cowrie  
  
git clone https://github.com/cowrie/cowrie.git  
  
cd cowrie  
  
python3 -m venv cowrie-env
```

```
source cowrie-env/bin/activate

pip install --upgrade pip

pip install -r requirements.txt

# copy template config and edit

cp etc/cowrie.cfg.dist etc/cowrie.cfg

# edit etc/cowrie.cfg - set hostname, enabled protocols, log targets (JSON)

# Example: configure journald/file output and enable JSON logging for downstream parsing
```

Create a systemd service (run as cowrie user) so it restarts on failure. Many guides show exact systemd unit files — follow Cowrie docs.

(Several step-by-step walk-throughs exist for exactly these steps.) codygula.com

Step F — Send logs to S3 & ELK

- Cowrie logs to var/log/cowrie and supports JSON output.
- Use awscli or a lightweight Filebeat agent to ship logs to:
 - S3 (for long-term storage and forensic retrieval).
 - Elastic/OpenSearch (for search/visualization).

Example to copy logs to S3 (cron or logrotate hook):

```
# install awscli and configure minimal IAM role on EC2 instead of local creds
```

```
sudo apt install awscli
```

```
# upload
```

```
aws s3 cp /home/cowrie/cowrie/log/ s3://my-honeypot-logs/cowrie/ --recursive
```

Better: attach an instance IAM role with s3:PutObject and use a small agent to push logs.

Step G — Enable VPC Flow Logs & GuardDuty

- In VPC console → Flow Logs → Create flow log for the honeypot subnet; destination S3 or CloudWatch Logs.
- Enable GuardDuty in the account and region — it will pick up suspicious behavior and integrate with CloudWatch events for alerting. [aws.github.io+1](https://aws.github.io/guardduty/)

Step H — Verify & test

- Check Cowrie UI/console logs.

- Run a simple external scan (from a separate test machine) against port 22 and see if it records the login attempts and commands.
 - Validate logs in S3 / ELK.
-

4) Scale / alternatives: T-Pot and MHN

- T-Pot: all-in-one, runs many honeypots + ELK Dashboards on one host (typically runs as an appliance image). Good for rapid deployment; many guides for EC2 exist. Use a larger instance (m5 / c5) for T-Pot. [InfoSec Write-ups+1](#)
 - MHN: deploy MHN server (one EC2) and many lightweight sensors (EC2 or on-prem) that register with the MHN server. MHN provides central deployment & data collection. Good for distributed honeypot fleets. See the MHN GitHub for sensor types and instructions. [GitHub](#)
-

5) Logging, analysis, and detection workflows

1. Collection:
 - Cowrie/T-Pot/MHN → JSON logs to S3 and to a central ELK/OpenSearch.
 - VPC Flow Logs → S3/CloudWatch → Athena queries for IP patterns.
 - CloudTrail & GuardDuty → CloudWatch Events (for correlation).
2. Processing:
 - Filebeat/Logstash → parse Cowrie JSON → index in Elasticsearch/OpenSearch.
 - Enrich logs with GeoIP, ASN, and known bad IP lists.
3. Analysis you should run regularly:
 - Top attacking IPs (count attempts).
 - Brute-force username/password pairs and common payloads.
 - Timeline of actions per IP (to detect multi-stage intrusions).
 - File artifacts (downloaded binaries) — quarantine and run in sandbox (Cuckoo / detonation lab).
 - Correlate with GuardDuty findings and CloudTrail to spot if any real resources were accessed.
4. Alerts:

- High volume of failed logins from same IP → alert.
- Successful session with file download → high severity (exfil/payload).
- Known C2 indicators observed in outgoing traffic from honeypot → alert and isolate.

(Plenty of guides show how to wire this into ELK or MHN dashboards.) [GitHub+1](#)

6) Data analysis example (what insights you can get)

- Attack surface profiling: which ports/protocols are most frequently probed, from which countries and ASNs.
- Attack behavior: common commands run (recon commands, privilege escalation attempts), scripts uploaded, file names used for malware.
- Malware capture: binaries and their metadata (hashes) for IOC creation and sandboxing.
- Trend analysis: time/day patterns, correlation with known exploit disclosures (e.g., new CVE scanning spikes).

Use Athena on S3 flow logs or Kibana dashboards to quickly answer such questions.

7) Risks, containment, and legal/ethical notes

1. **Pivot / abuse risk:** a fully interactive honeypot can be used by attackers to launch attacks. To reduce risk:
 - Strongly restrict outbound traffic from honeypot (prevent direct internet access except to controlled update hosts).
 - Monitor and block any attempt to escalate or access AWS metadata endpoints (IMDS). Block IMDS access from honeypot. (IMDS is a common AWS pivot target.)
 - Use network ACLs and routing rules to isolate the honeypot.
2. **Legal & policy:** running a honeypot that captures personal data or allows law-breaking activity may have legal implications. Check your organization's policy and, if needed, legal counsel.
3. **Do NOT connect honeypots to production networks** — always isolate. Multiple tutorials and AWS posts stress isolating honeypots and using monitoring services.
[AWS Documentation+1](#)

8) Cost & cleanup

- Cost drivers: EC2 instance hours (especially bigger types for T-Pot/MHN), S3 storage for logs, data transfer if you ship logs out of region, EBS volumes, and any ELK/OpenSearch provisioning.
 - Test with t3.micro first; scale up for heavier logging or T-Pot.
 - Cleanup: terminate EC2, delete EBS snapshots, remove S3 objects, disable VPC Flow Logs, and remove IAM roles/policies.
-

9) Quick checklist / runbook (copyable)

- Create separate AWS account or VPC for honeypots.
 - Create VPC, public subnet, internet gateway.
 - Create security groups (admin restricted, honeypot ports open).
 - Launch EC2 (Ubuntu), attach IAM role for S3 PutObject and CloudWatch.
 - Install Cowrie (or T-Pot/MHN), configure to output JSON.
 - Configure VPC Flow Logs → S3.
 - Configure CloudWatch / GuardDuty / CloudTrail.
 - Configure log forwarding to ELK/OpenSearch and set dashboards.
 - Set automated snapshot & retention policy for S3 logs.
 - Test by running controlled scans and verify logs & alerts.
 - Hardening: block IMDS for the honeypot (e.g., iptables or VPC route policies), restrict outbound traffic.
 - Document cleanup steps & cost thresholds.
-

10) Resources & further reading (selected)

- AWS VPC security best practices (VPC Flow Logs + GuardDuty recommendations).
[AWS Documentation](#)
- Cowrie SSH honeypot install guides (practical walkthroughs). [codygula.com](#)
- Modern Honey Network (MHN) GitHub (centralized management for many sensors).
[GitHub](#)

- T-Pot / all-in-one honeypot platform guides and deployments. [InfoSec Write-ups](#)
- AWS Security blog posts about deception/honeypot approaches on AWS. [Amazon Web Services, Inc.](#)